

# forest fire prediction

November 2, 2022

## 1 data Details

- Machine Learning LifeCycle
  1. Data Ingestion
  2. EDA
  3. Preprocessing
  4. Model Building
  5. Performance Metrics
- Problem Statement
  - The dataset includes 244 instances that regroup a data of two regions of Algeria,namely 122 instances for each region.  
The period from June 2012 to September 2012.  
The dataset includes 11 attribues and 1 output attribue (class)  
The 244 instances have been classified into fire (138 classes) and not fire (106 classes)
- Feature Description
  1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012) Weather data observations
  2. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
  3. RH : Relative Humidity in %: 21 to 90
  4. Ws :Wind speed in km/h: 6 to 29
  5. Rain: total day in mm: 0 to 16.8 FWI Components
  6. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
  7. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
  8. Drought Code (DC) index from the FWI system: 7 to 220.4
  9. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
  10. Buildup Index (BUI) index from the FWI system: 1.1 to 68
  11. Fire Weather Index (FWI) Index: 0 to 31.1
  12. Classes: two classes, namely Fire and not Fire
- Machine Learning Models Used
  1. Linear Regression
  2. Ridge Regression
  3. Lasso Regression
  4. Elastic-Net Regression

- Assumption of Linear Regression <https://towardsdatascience.com/assumptions-of-linear-regression-fdb71ebeaa8b>

```
[2]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
[3]: df = pd.read_csv('Algerian_forest_fires_dataset_UPDATE.csv',header=1)
```

## 2 Data Cleaning

```
[4]: df
```

```
[4]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	\
0	01	06	2012	29	57	18	0	65.7	3.4	7.6	1.3	3.4	
1	02	06	2012	29	61	13	1.3	64.4	4.1	7.6	1	3.9	
2	03	06	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	
3	04	06	2012	25	89	13	2.5	28.6	1.3	6.9	0	1.7	
4	05	06	2012	27	77	16	0	64.8	3	14.2	1.2	3.9	
..	..	...	...	...	..	..	...	...	...	...	...	...	
241	26	09	2012	30	65	14	0	85.4	16	44.5	4.5	16.9	
242	27	09	2012	28	87	15	4.4	41.1	6.5	8	0.1	6.2	
243	28	09	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	
244	29	09	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	
245	30	09	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	
	FWI		Classes										
0	0.5		not fire										
1	0.4		not fire										
2	0.1		not fire										
3	0		not fire										
4	0.5		not fire										
..	...		...										
241	6.5		fire										
242	0		not fire										
243	0.2		not fire										
244	0.7		not fire										
245	0.5		not fire										

```
[246 rows x 14 columns]
```

```
[5]: # how big the Data is
df.shape
```

```
[5]: (246, 14)
```

```
[6]: # how the data is distributed
df.describe().T
```

```
[6]:
```

	count	unique	top	freq
day	246	33	01	8
month	245	5	07	62
year	245	2	2012	244
Temperature	245	20	35	29
RH	245	63	64	10
Ws	245	19	14	43
Rain	245	40	0	133
FFMC	245	174	88.9	8
DMC	245	167	7.9	5
DC	245	199	8	5
ISI	245	107	1.1	8
BUI	245	175	3	5
FWI	245	128	0.4	12
Classes	244	9	fire	131

```
[7]: df.columns = df.columns.str.strip()
```

```
[8]: df[df['day']=='Sidi-Bel Abbes Region Dataset']['day'] = np.nan
df[df['day']=='day']['day']= np.nan
df[df['year']=='year']['year']= np.nan
df[df['day']=='Sidi-Bel Abbes Region Dataset']
df.iloc[122,:][0] = np.nan
df1 = df.drop(123)
```

```
[9]: df1.dropna(inplace=True)
```

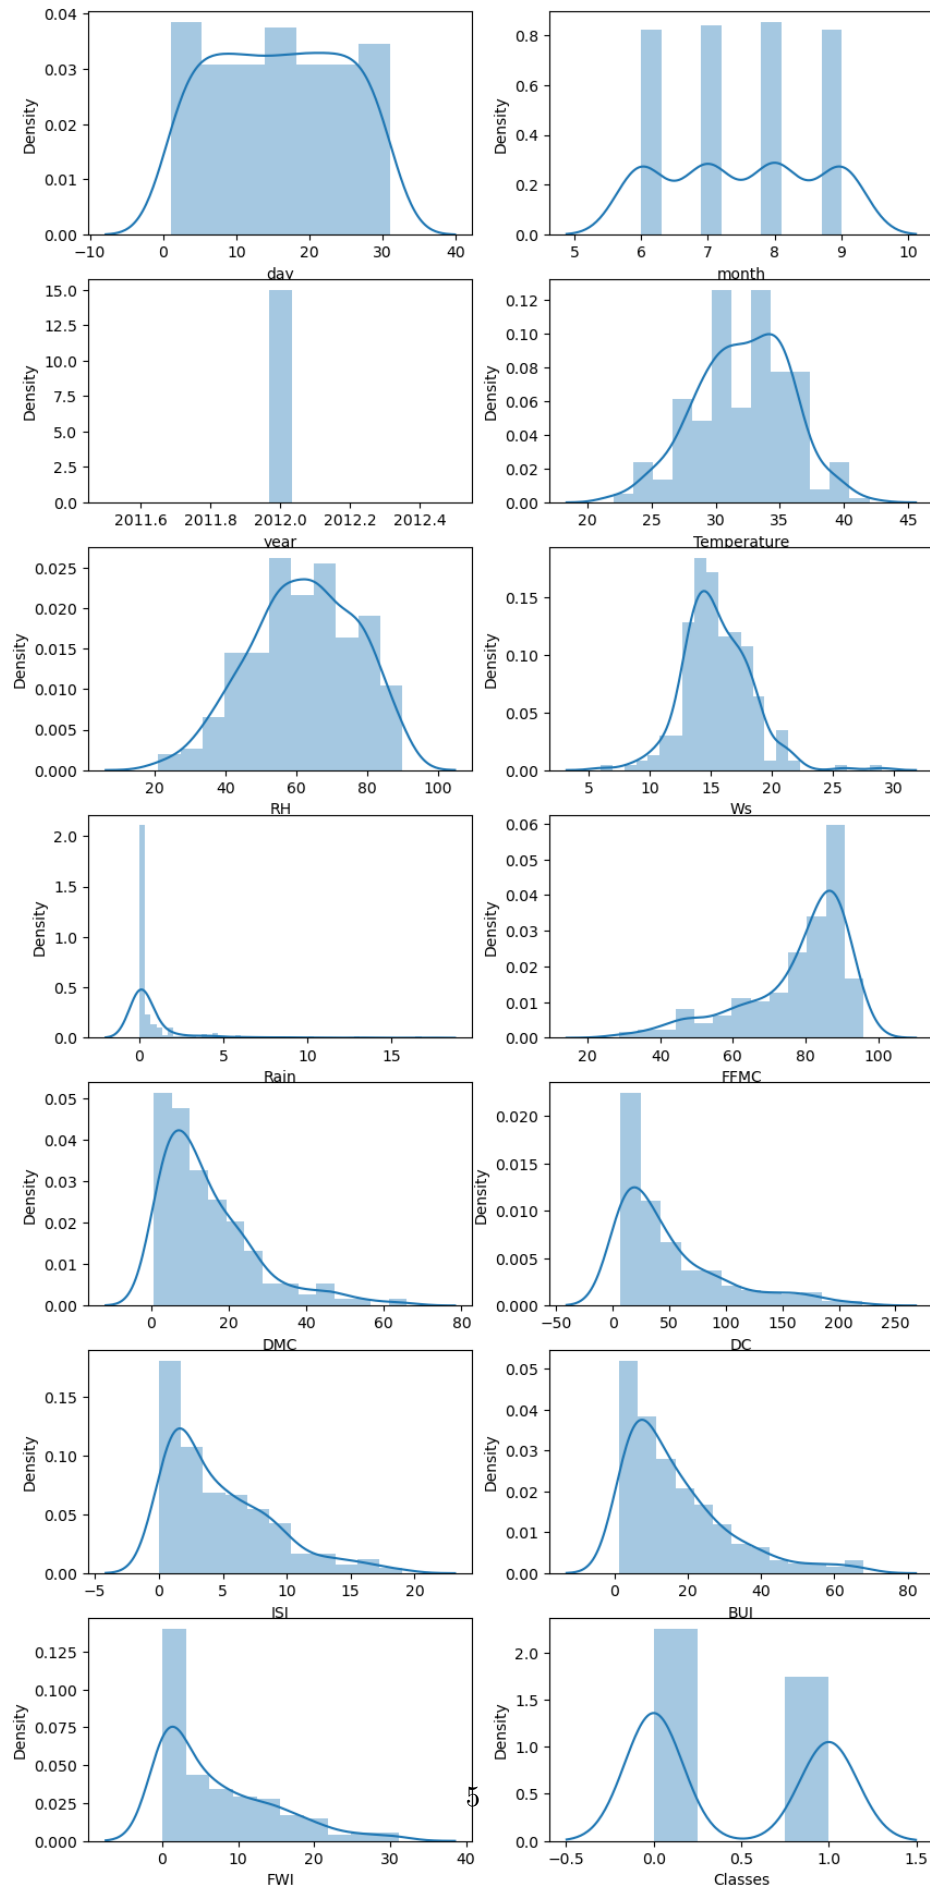
```
[10]: df1['Classes'] = df1['Classes'].str.strip()
```

```
[11]: cl = list(df1.columns)
for i in cl[:-1]:
    df1[i] = pd.to_numeric(df1[i])
```

```
[12]: from sklearn.preprocessing import LabelEncoder
lb = LabelEncoder()
df1['Classes'] = lb.fit_transform(df1['Classes'])
```

```
[13]: plt.figure(figsize=(10,25))
for i in zip(df1.columns,list(range(len(df1.columns)))):
```

```
plt.subplot(8,2,i[1]+1)  
sns.distplot(df1[i[0]])
```



### 3 Model Buliding

```
[14]: from sklearn.linear_model import LinearRegression,Ridge,Lasso,ElasticNet
      from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
      from sklearn.model_selection import cross_val_score
```

```
[15]: x = df1.drop('Temperature',axis=1)
      y = df1['Temperature']
      from sklearn.model_selection import train_test_split
      X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.
      ↪2,random_state=32)
```

### 4 Linear Regression:

```
[16]: lr = LinearRegression()
      lr.fit(X_train,y_train)
      lr_predict = lr.predict(X_test)
      print("Train Score:",lr.score(X_train,y_train))
      print('Test Score: ',lr.score(X_test,y_test))
      from sklearn.metrics import r2_score
      score=r2_score(y_test,lr_predict)
      print("The R2 Score for the model builded is",score)
      Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
      print("The Adjusted R Square for the model is",Adjusted_r)
```

Train Score: 0.5796076040993602

Test Score: 0.5761842043675531

The R2 Score for the model builded is 0.5761842043675531

The Adjusted R Square for the model is 0.41876690884693

```
[17]: from sklearn.preprocessing import StandardScaler
      standard = StandardScaler()
      X_train_s = standard.fit_transform(X_train)
      X_test_s =standard.transform(X_test)
```

```
[18]: # After Standartization
      lr.fit(X_train_s,y_train)
      lr_predict = lr.predict(X_test_s)
      print("Train Score:",lr.score(X_train_s,y_train))
      print('Test Score: ',lr.score(X_test_s,y_test))
      from sklearn.metrics import r2_score
      score=r2_score(y_test,lr_predict)
      print("The R2 Score for the model builded is",score)
```

```
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test_s.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```


Train Score: 0.5796076040993601

Test Score: 0.5761842043675535

The R2 Score for the model builded is 0.5761842043675535


The Adjusted R Square for the model is 0.41876690884693046

```
[51]: # ASSUMPTION
plt.scatter(y_test,lr_predict)
plt.xlabel("Test Truth ")
plt.ylabel("Test Predicted ")
plt.show()
```



```
[42]: residuals=y_test-lr_predict
sns.displot(residuals,kind='kde')
```


```
[42]: <seaborn.axisgrid.FacetGrid at 0x7f2481cb6410>
```



```
[43]: plt.scatter(lr_predict,residuals)
```

```
[43]: <matplotlib.collections.PathCollection at 0x7f2481afd0c0>
```






```
[46]: from statsmodels.stats.outliers_influence import variance_inflation_factor
      {x.columns[i]: variance_inflation_factor(x.values, i) for i in range(1, x.
      ↪shape[1])}
```

```
[46]: {'month': 1.0593677394567451,
      'year': 381.03887723157646,
      'RH': 2.860679742317785,
      'Ws': 1.2618825859478544,
      'Rain': 1.5751316142697394,
      'FFMC': 5.301788307992688,
      'DMC': 77.36193951221615,
      'DC': 25.199194461781985,
      'ISI': 23.250581619252316,
      'BUI': 172.04514981024676,
      'FWI': 41.1884258115508,
      'Classes': 3.474838521935455}
```

```
[50]: plt.scatter(y_test, residuals)
      plt.plot(y_test, [0]*len(y_test))
```

```
[50]: [<matplotlib.lines.Line2D at 0x7f24742995a0>]
```



## 5 Ridge Regression

[ ]:

```
[19]: rid = Ridge()
print(rid.fit(X_train_s,y_train))
lr_predict = rid.predict(X_test_s)
print("Train Score:",rid.score(X_train_s,y_train))
print('Test Score: ',rid.score(X_test_s,y_test))
from sklearn.metrics import r2_score
score=r2_score(y_test,lr_predict)
print("The R2 Score for the model builded is",score)
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test_s.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)
```

```
Ridge()
Train Score: 0.5791326452390795
Test Score: 0.5916874165048187
The R2 Score for the model builded is 0.5916874165048187
The Adjusted R Square for the model is 0.4400284569208942
```

```
[21]: # Lasso Regression
las =Lasso(alpha=0.04,max_iter=1000)
print(las.fit(X_train_s,y_train))
lr_predict = las.predict(X_test_s)
print("Train Score:",las.score(X_train_s,y_train))
print('Test Score: ',las.score(X_test_s,y_test))
from sklearn.metrics import r2_score
score=r2_score(y_test,lr_predict)
print("The R2 Score for the model builded is",score)
Adjusted_r=1-(1-score)*(len(y_test)-1)/(len(y_test)-X_test_s.shape[1]-1)
print("The Adjusted R Square for the model is",Adjusted_r)

a = np.mean(cross_val_score(estimator=las,X = x,y = y,cv=10,scoring='r2'))
print(a)
```

```
Lasso(alpha=0.04)
Train Score: 0.5758497343585591
Test Score: 0.6060660158722289
The R2 Score for the model builded is 0.6060660158722289
The Adjusted R Square for the model is 0.4597476789104853
0.036977819359694984
```

```
[22]: # Hyper Parameter Tunning
from sklearn.model_selection import GridSearchCV
para = {}
para['alpha'] = list(np.linspace(0,0.2,40))
para['max_iter'] = [100,10000,10000]
grit = GridSearchCV(estimator=las,param_grid=para,verbose=1,cv = 10)
grit.fit(X_train_s,y_train)
grit.best_score_
```

Fitting 10 folds for each of 120 candidates, totalling 1200 fits

```
[22]: 0.5053909597293631
```

```
[39]: # Ridge Hyperparameter tuning
from sklearn.model_selection import GridSearchCV
para = {}
para['solver'] =['auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', 'saga', 'lbfgs']
para['alpha'] = list(np.linspace(0,300,100))
para['max_iter'] = [10000,100000]
grit = GridSearchCV(estimator=rid,param_grid=para,verbose=1,cv = 10,n_jobs=43)
grit.fit(X_train_s,y_train)
grit.best_score_
```

Fitting 10 folds for each of 1600 candidates, totalling 16000 fits

```
[39]: 0.5128037686657727
```

```
[40]: grit.best_params_
```

```
[40]: {'alpha': 45.45454545454545, 'max_iter': 10000, 'solver': 'sparse_cg'}
```

```
[ ]:
```